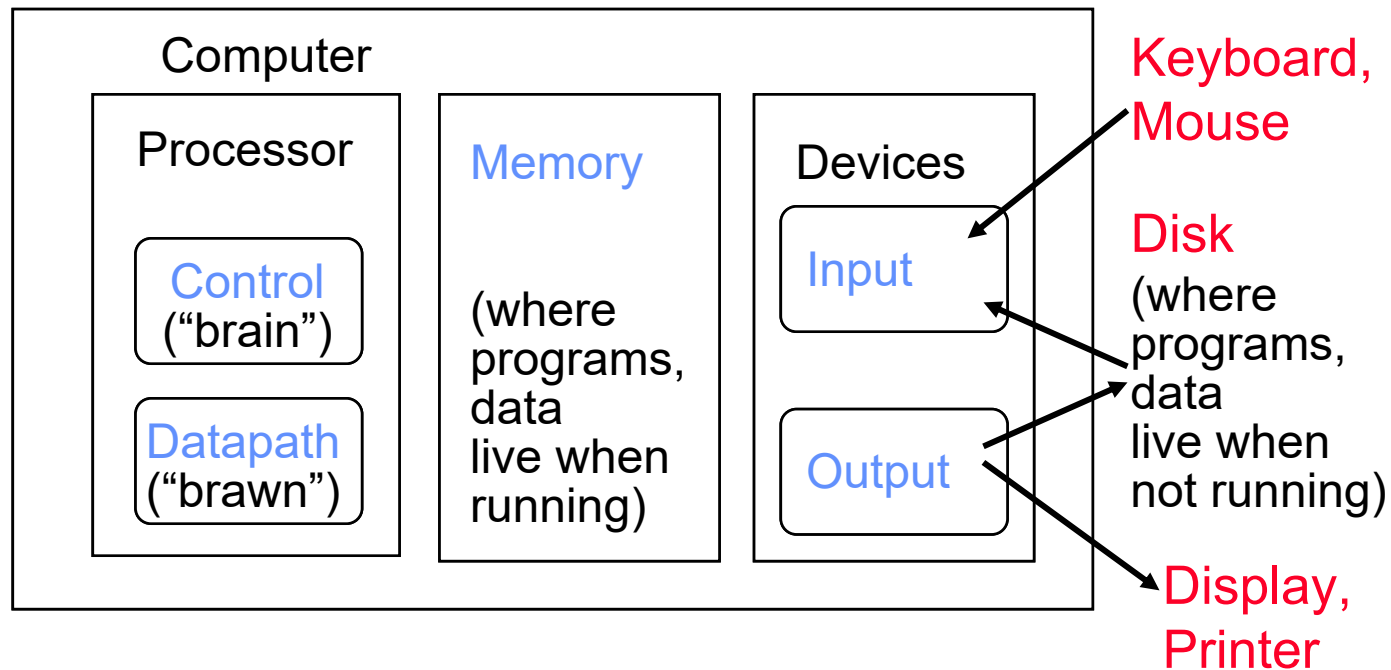



Five Components of a Computer


- Any computer, no matter how primitive or advanced, can be divided into five parts:




First Example: iPad




Second Example: x86 Motherboard



Third Example: Cray XT6 Supercomputer Blade



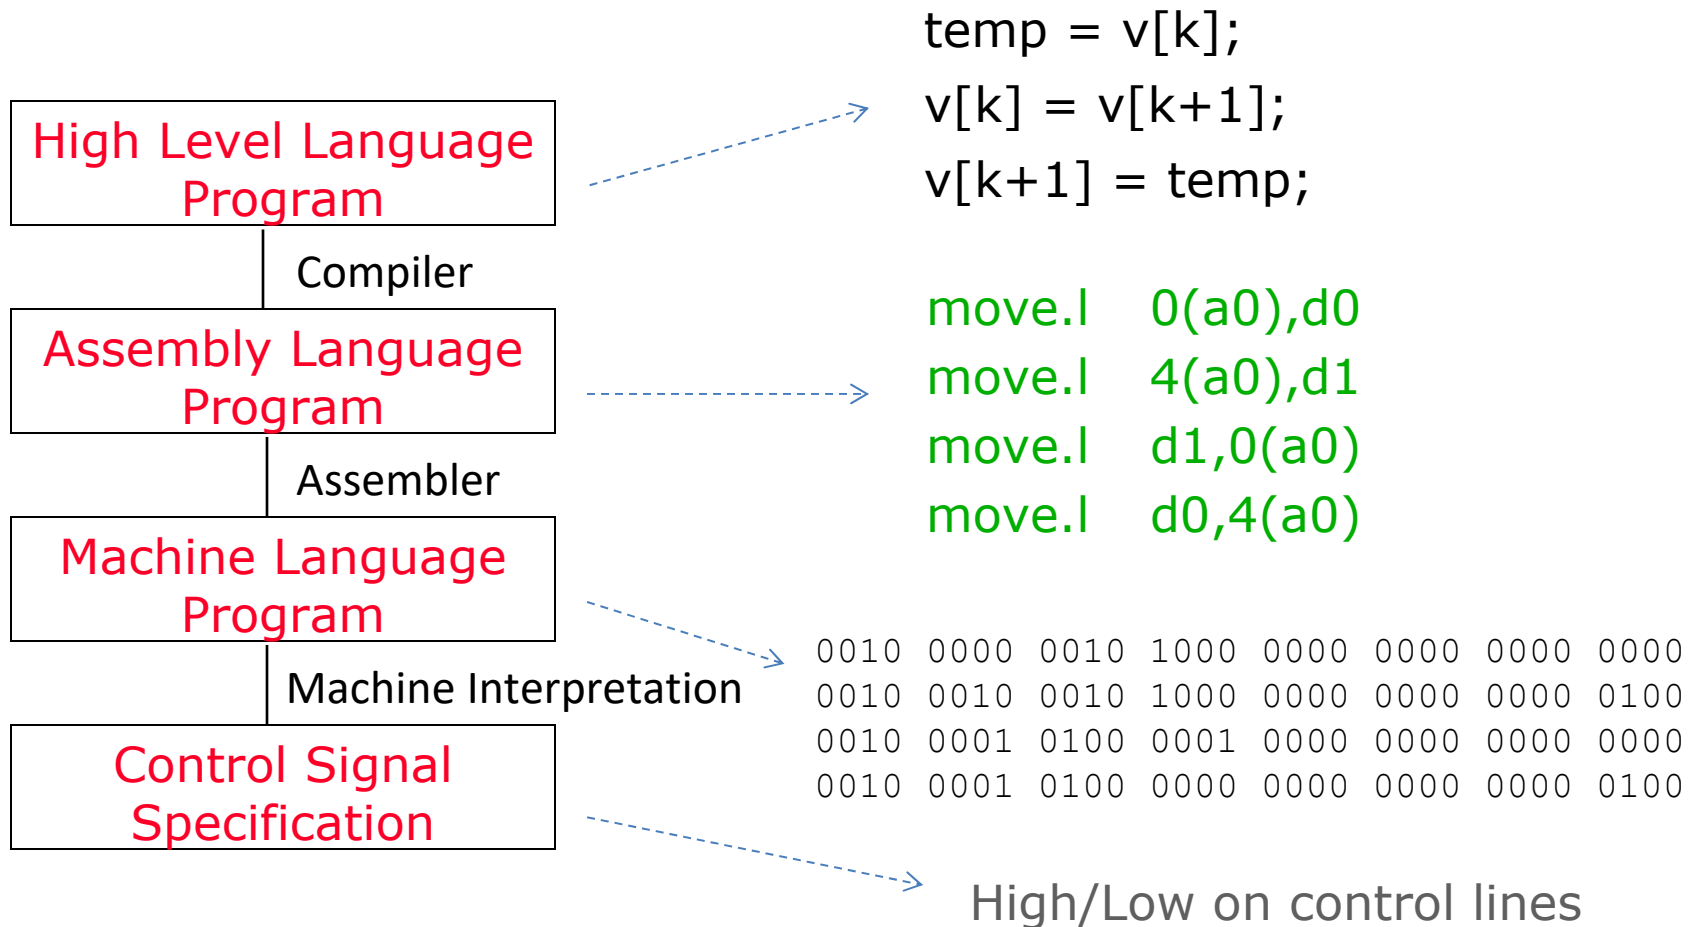
Fourth Example: 68000 Mini-board (68KMB)



Computers Are Pretty Simple – Huh?

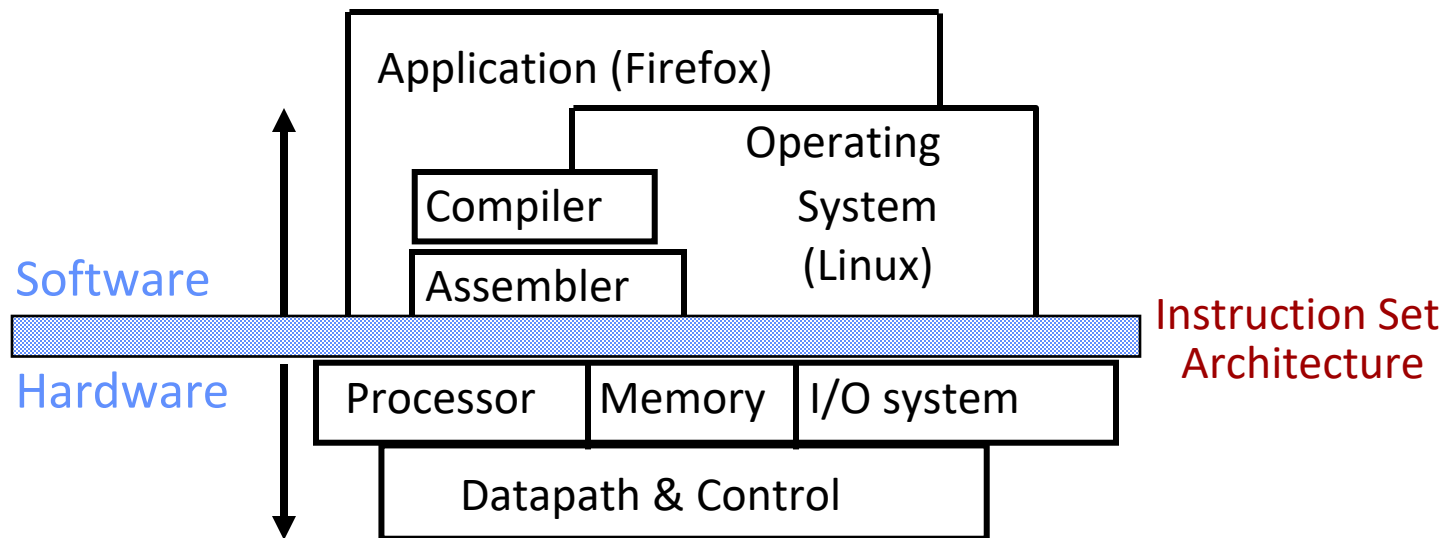
- Computers only work with binary signals
 - Signal on a wire is either 0 or 1
 - Usually called a “bit”
 - More complex stuff (numbers, characters, strings, pictures)
 - Must be build from multiple bits
- Built out of simple logic gates that perform Boolean logic
 - AND, OR, NOT, Adders, Comparators, etc.
- And memory cells that preserve bits over time
 - Flip-flops, registers, SRAM cells, DRAM cells, Flash, etc.
- To get hardware to do anything, need to break it down to bits
 - Strings of bits that tell the hardware what to do are called
 - instructions
 - A sequence of instructions is called
 - machine code or machine-language program

Running An Application



Instruction-Set Architecture (ISA)

- Instruction-Set Architecture
 - Important abstraction
 - Interface between hardware and lowest-level software
 - Portion of computer visible to programmer
 - Assembly language, programmer's model, instructions/format




Instruction-Set Architecture (ISA)

Examples

- ARM, MIPS, Intel IA32 (x86), Sun SPARC, PowerPC, IBM 390, Intel IA64, M68000, etc.
 - These are all ISAs
- Many different implementations can implement the same ISA (family)
 - 8086, 386, 486, Pentium, Pentium II, Pentium IV implement IA32
 - Of course they continue to extend it, while maintaining binary compatibility
- ISAs last a long time
 - x86 has been in use since the 70s
 - IBM 390 started as IBM 360 in 60s
 - Stable interface between software and hardware
- Micro-architecture
 - Processor design techniques used to implement the ISA

A First Look at Memory


- Memory holds both instructions and data
- Logically organized as an array of **locations**, each with an **address**



A First Look at Memory

- Memory holds both instructions and data
- Logically organized as an array of **locations**, each with an **address**

- Process
 - Running program + State
 - Each process “sees” full address space due to an abstraction, called virtual memory
- Compiler + run-time system
 - Control memory allocation



Word Size of Processor

- Nominal size of integer-valued data
 - Includes addresses used to index memory
- Until recently, most machines used 32-bit (4-byte) words
 - Limits addresses to 4 Gigabytes
 - Become too small for memory-intensive applications
- Most current machines use 64-bit (8-byte) words
 - Potential address space: $2^{64} \approx 1.8 \times 10^{19}$ bytes (16 exabytes)
- For backward-compatibility many processors support different word sizes
 - Always a power-of-2 in the number of bytes: 1,2,3,4,5,6, ...

C Data Types

- Typical size of C objects in bytes

Data Object	32-bit Machine	64-bit Machine
char	1 byte	1 byte
int	4 bytes	4 bytes
short int	2 bytes	2 bytes
long int	4 bytes	8 bytes
long long	8 bytes	8 bytes
float	4 bytes	4 bytes
double	8 bytes	8 bytes
long double	8 bytes	16 bytes
pointer	4 bytes	8 bytes

Word-Oriented View of Memory

- Addresses specify locations of bytes in memory
 - Word address is address of first byte in word
 - Address of successive words differ by 4 (32-bit) or 8 (64-bit)

				0x01C
				0x018
				0x014
				0x010
Byte 15	Byte 14	Byte 13	Byte 12	0x00C
				0x008
				0x004
Byte 3	Byte 2	Byte 1	Byte 0	0x000

K=32-bits

Byte 15	Byte 14	0x00E
Byte 13	Byte 12	0x00C
		0x00A
		0x008
		0x006
		0x004
Byte 3	Byte 2	0x002
Byte 1	Byte 0	0x000

K=16-bits


Byte 15	0x00F
Byte 14	0x00E
Byte 13	0x00D
Byte 12	0x00C
	0x00B
	0x00A
	0x009
	0x008
	0x007
	0x006
	0x005
	0x004
Byte 3	0x003
Byte 2	0x002
Byte 1	0x001
Byte 0	0x000

K=8-bits

Variable Declarations and Memory

- Variable declarations
 - `int a, b;`
 - find two locations in memory in which to store two integers (1 word each)

- Memory Layout
 - Assume word size of 32-bits



Variable Assignments and Memory

- Assignment (LHS = RHS)
 - `int a = 0, b = -1;`
 - LHS must evaluate to a memory location (a variable)
 - RHS must evaluate to a value (possibly an address)

				0x01C
				0x018
				0x014
FF	FF	FF	FF	0x010 b
				0x00C
				0x008
00	00	00	00	0x004 a
				0x000

Variable Assignments and Memory


- Assignment (LHS = RHS)
 - `int a = 0, b = -1;`
 - LHS must evaluate to a memory location (a variable)
 - RHS must evaluate to a value (possibly an address)

- Example
 - `a = a + b;`

				0x01C
				0x018
				0x014
FF	FF	FF	FF	0x010 b
				0x00C
				0x008
00	00	00	00	0x004 a
				0x000


Pointers and Memory

- Pointer declaration
 - `int *a;`
 - declares a variable `a` that is a pointer to an integer data item
- Pointer assignment
 - `a = &b;`
 - assigns `a` the address where `b` is stored



Pointers and Memory

- Pointer dereferencing
 - `*a;`
 - the value at the memory address given by the value of `*a`




Pointers and Memory

- Pointer dereferencing

- $*a$;
- the value at the memory address given by the value of $*a$


- Example

- $b = *a + 1$;




Pointers and Memory

- Pointer arithmetic
 - $a = a + 1;$
 - the value of the pointer is incremented to point at the next data object (integer) in memory



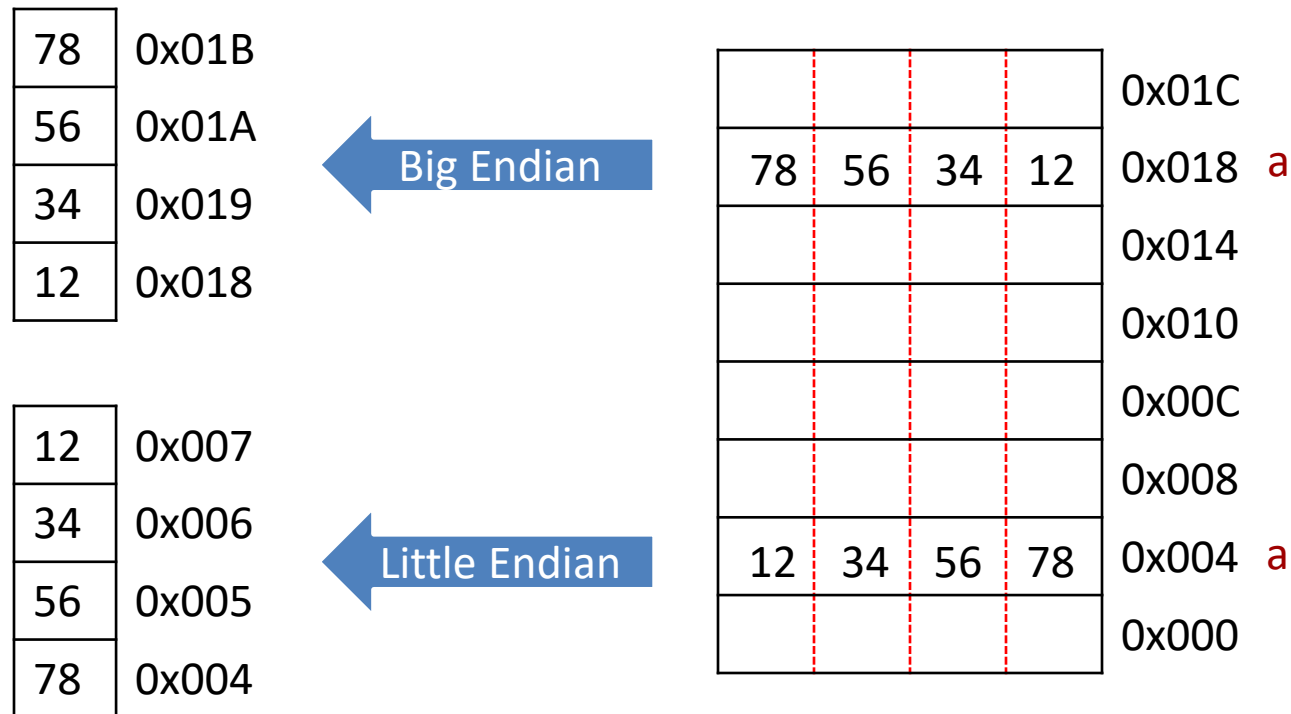
Pointers and Memory

- Pointer arithmetic
 - $a = a + 1;$
 - the value of the pointer is incremented to point at the next data object (integer) in memory



Byte Ordering

- How do we order the bytes in a word?
 - `int a = 0x12345678;`



Endianness: Big or Little?

- Big Endian
 - Address of most-significant byte = (lowest) address of word
 - IBM 360, M68000, MIPS, SPARC
- Little Endian
 - Address of least-significant byte = (lowest) address of word
 - Intel x86, ARM, DEC Vax & Alpha

When Endianness Matters

- When you
 - store words then load bytes

```
#include <stdio.h>

int main() {
    int fourBytes = 0x12345678;
    char oneByte = *(char*) &fourBytes;
    if (oneByte == 0x78)
        puts("little endian");
    else
        puts("big endian");

    return 0;
}
```

Big Endian

78	0x007
56	0x006
34	0x005
12	0x004

fourBytes

Little Endian

12	0x007
34	0x006
56	0x005
78	0x004

fourBytes

When Endianness Matters

- When you
 - share data between machines with different endianness
 - When moving between 32-bit and 64-bit word sizes

12	0x007
34	0x006
56	0x005
78	0x004

x86-32 (Little Endian)

00	0x00B
00	0x00A
00	0x009
00	0x008
12	0x007
34	0x006
56	0x005
78	0x004

X86-64 (Little Endian)

78	0x007
56	0x006
34	0x005
12	0x004


32-bit (Big Endian)

78	0x00B
56	0x00A
34	0x009
12	0x008
00	0x007
00	0x006
00	0x005
00	0x004


64-bit (Big Endian)

Alignment Issues

- Some ISAs require words to be at an even addresses
 - Consider the following word (2-byte) access




- Consider the following word (4-byte) access



C Example

What is the size of this structure?

```
struct foo {  
    char    sm;      /* 1 byte */  
    short   med;      /* 2 bytes */  
    char    sm1;     /* 1 byte */  
    int     lrg;      /* 4 bytes */  
}
```



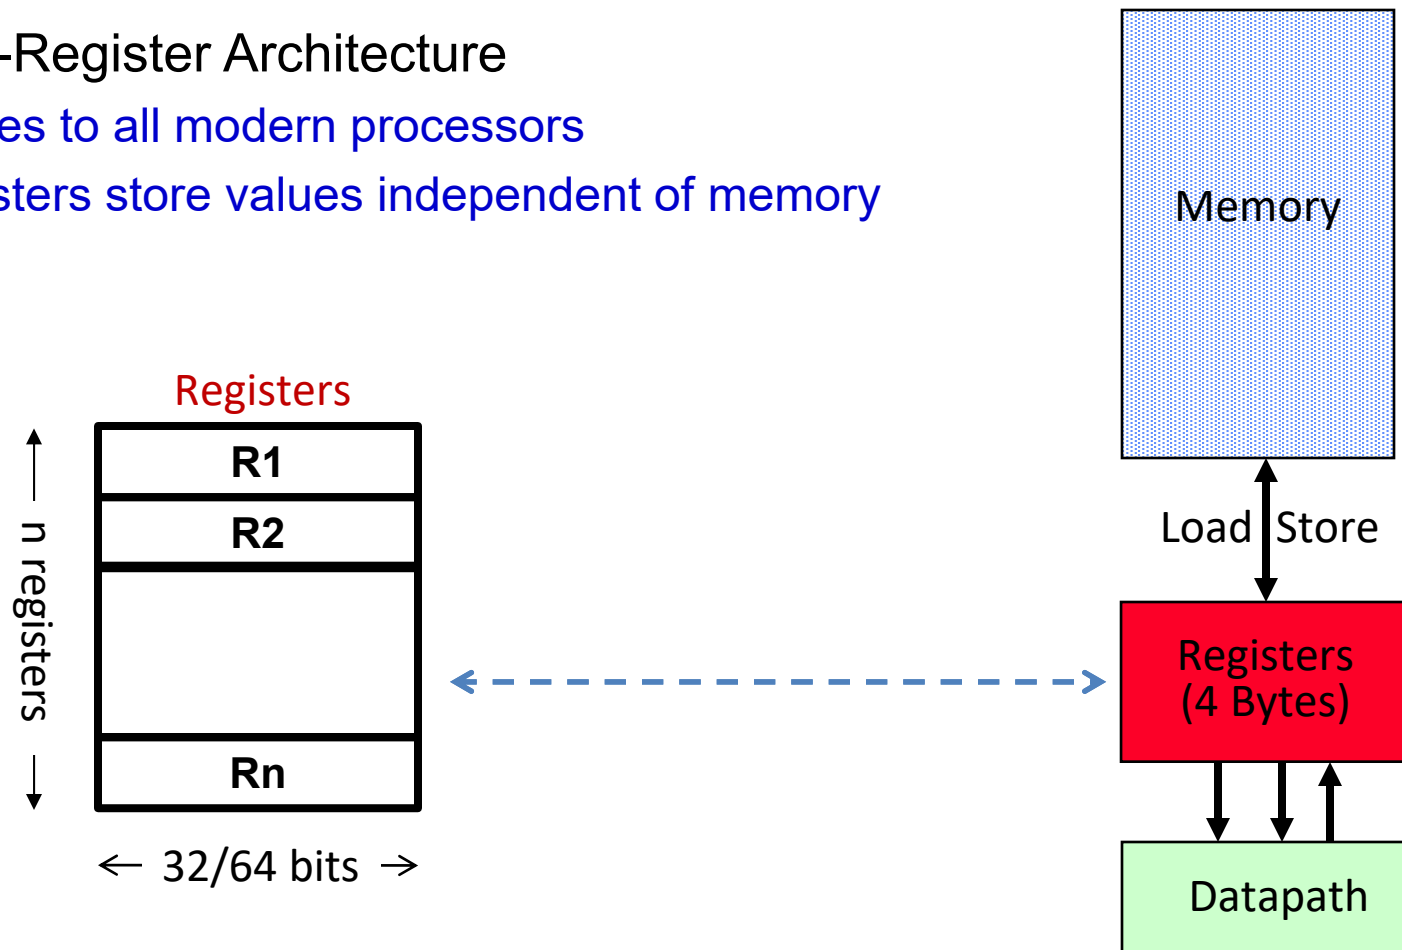
- Historically
 - Early machines (IBM 360, circa 1964) required alignment
 - Removed in the 1970s to reduce impact on programmers (IBM 370, x86)
 - Reintroduced by RISC to improve performance
 - Removed by some RISCs to simplify software

Today

- Major Components of Computer
- Instruction-Set Architecture (ISA)
- Memory
- **Registers**
- Instruction Cycle

A First Look at Registers

- General-Register Architecture
 - Applies to all modern processors
 - Registers store values independent of memory

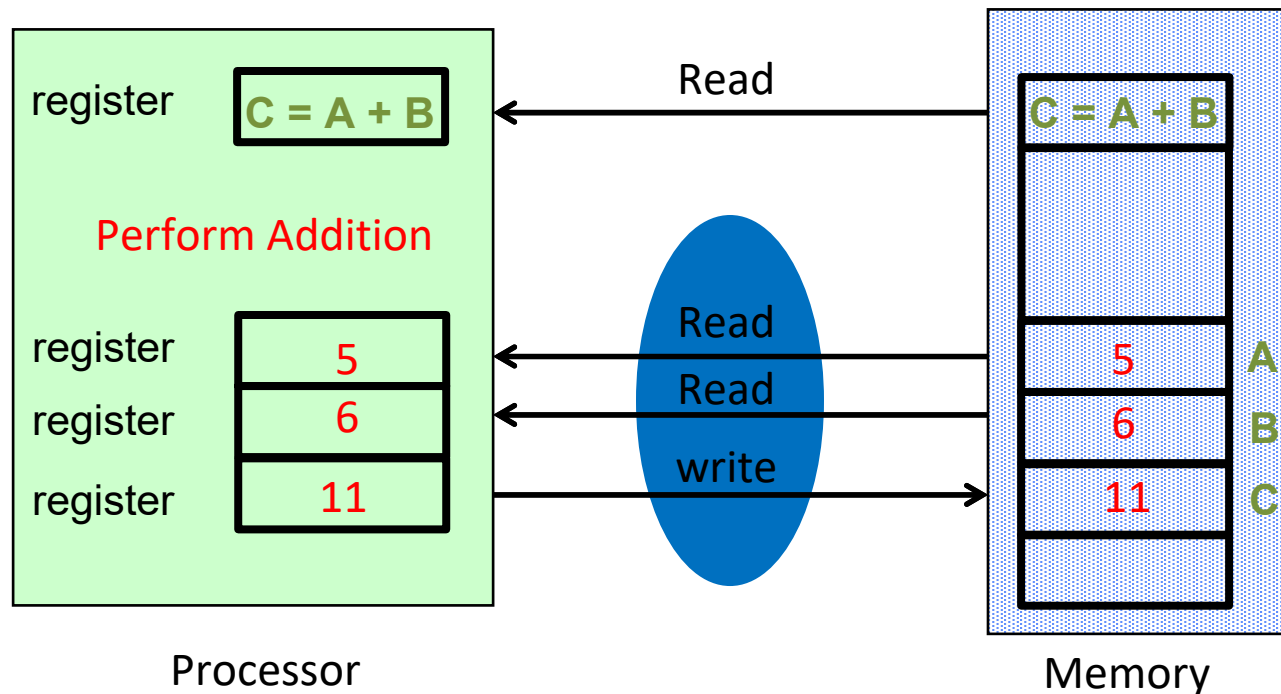


Today

- Major Components of Computer
- Instruction-Set Architecture (ISA)
- Memory
- Registers
- **Instruction Cycle**

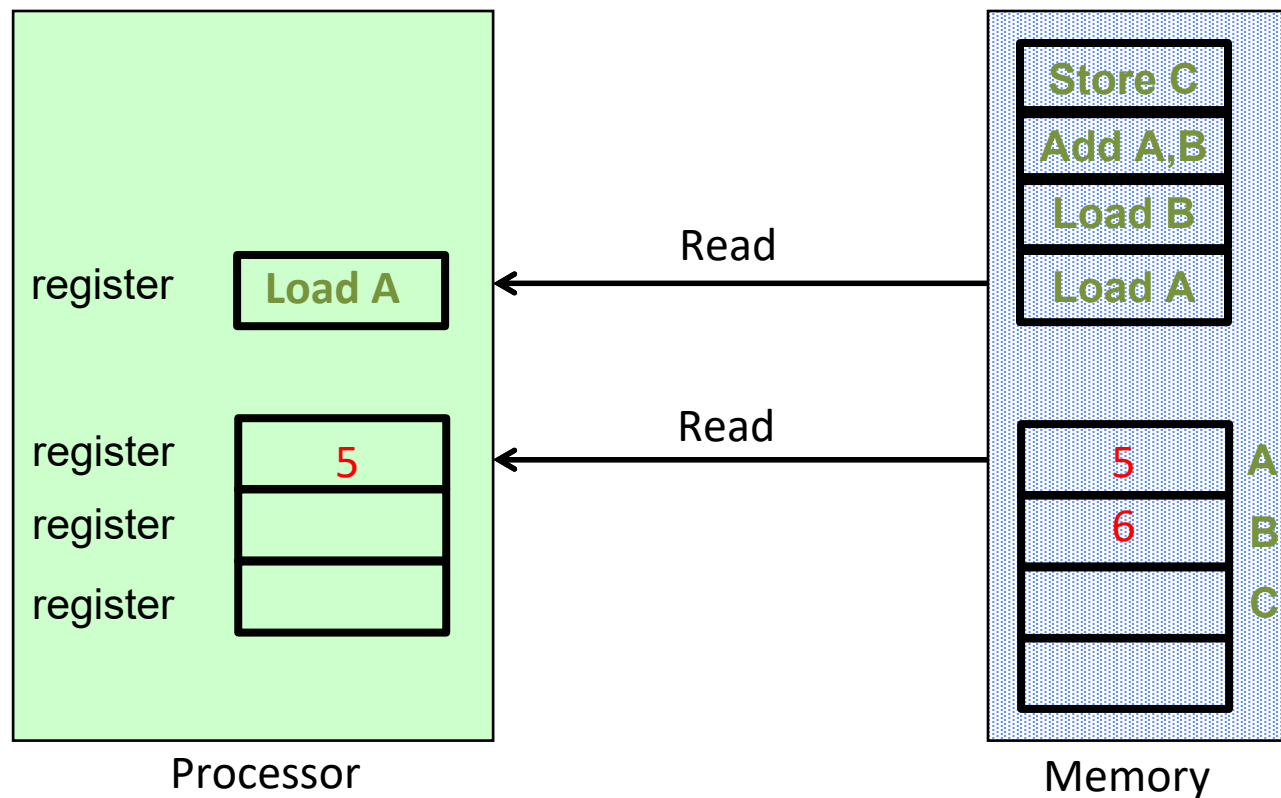
A First Look at the Processor

- Responsible for reading program instructions from memory and executing them



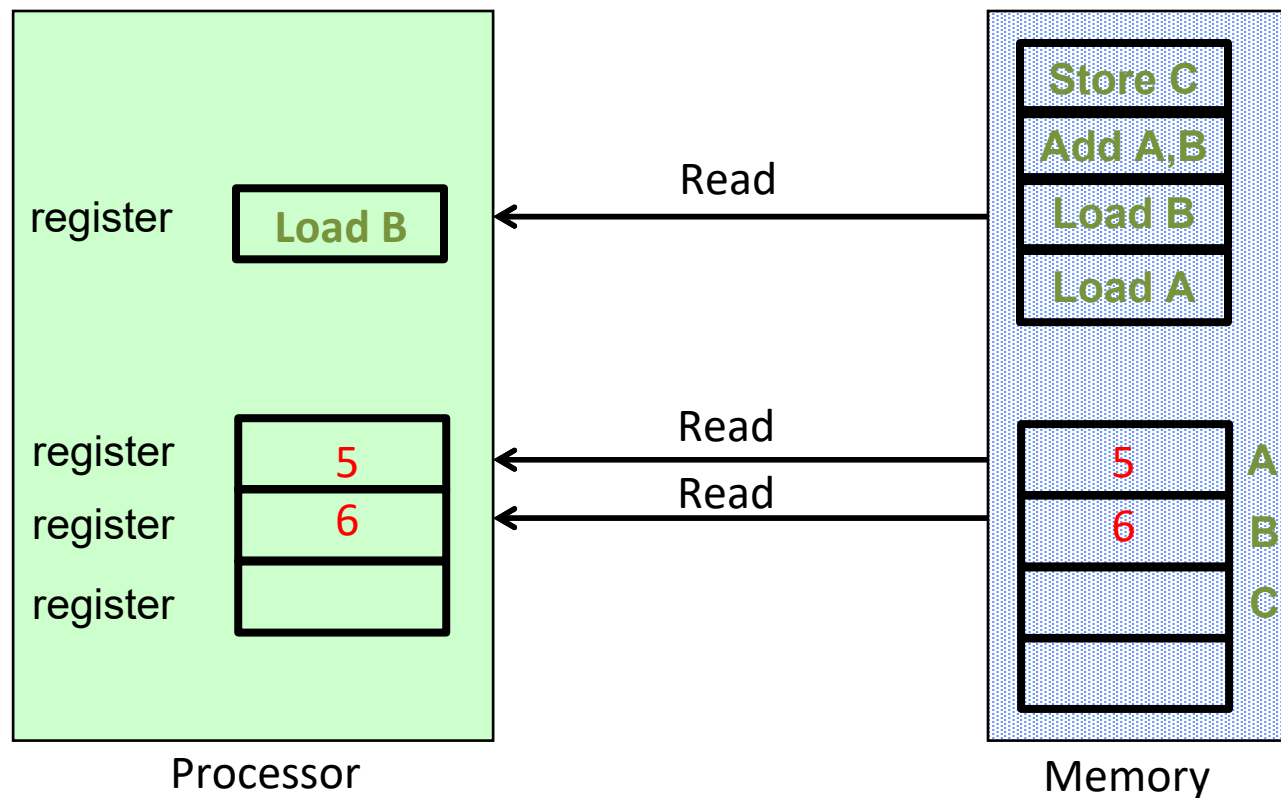
Instruction Granularity

- A statement such as $C = A + B$ is implemented as a sequence of primitive machine instructions that access memory at most once



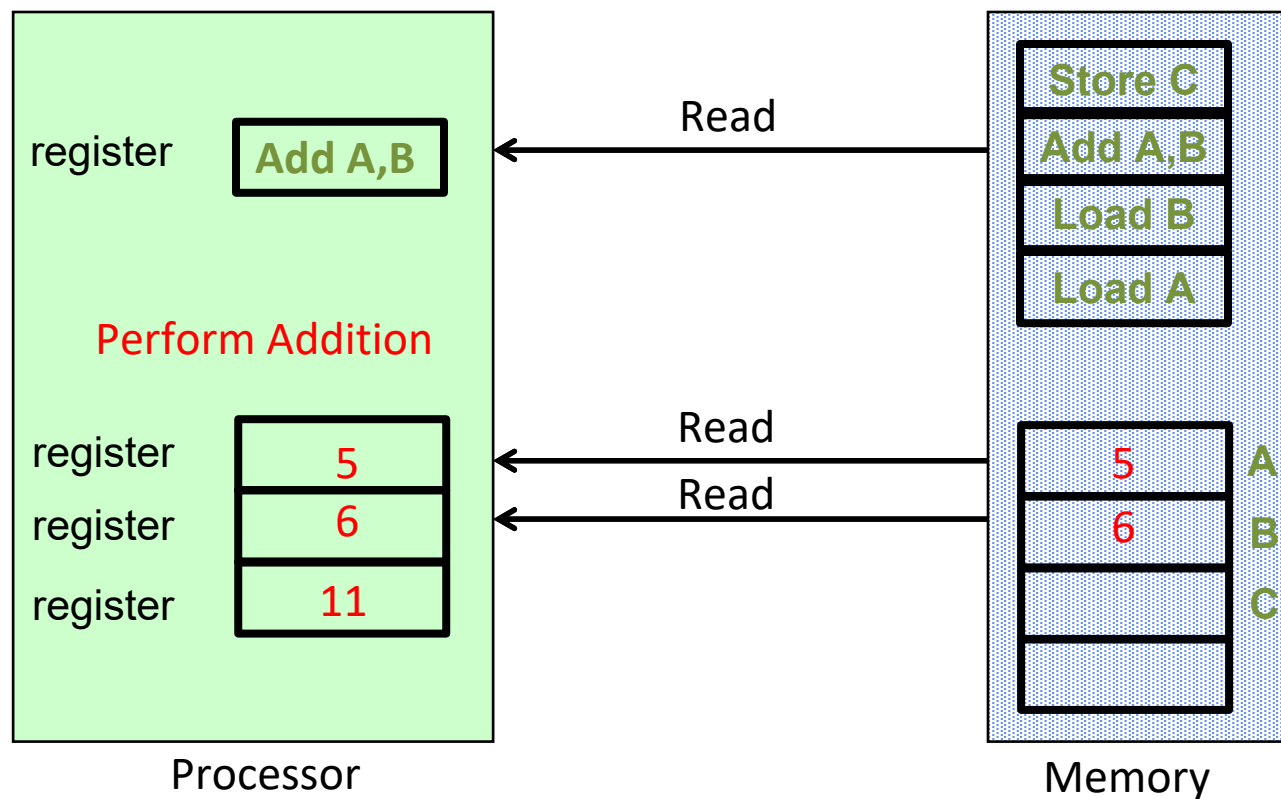
Instruction Granularity

- A statement such as $C = A + B$ is implemented as a sequence of primitive machine instructions that access memory at most once



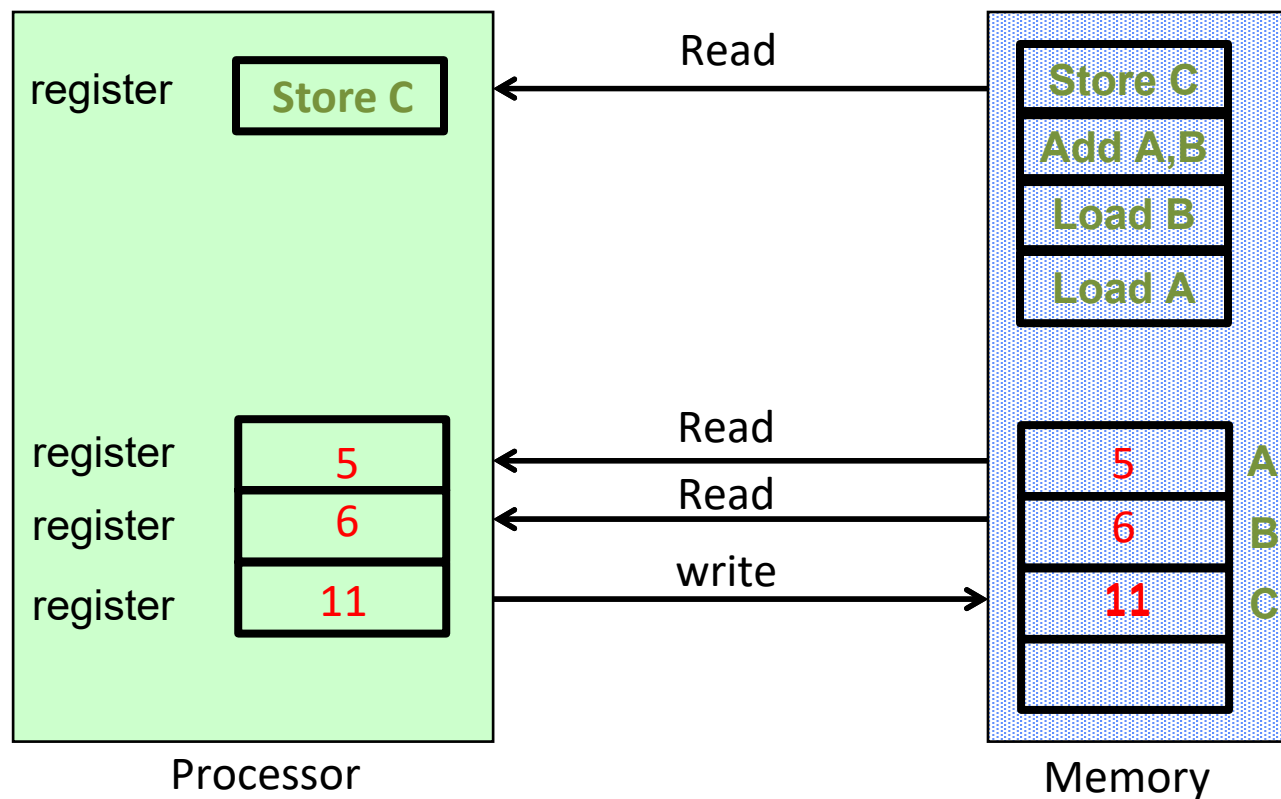
Instruction Granularity

- A statement such as $C = A + B$ is implemented as a sequence of primitive machine instructions that access memory at most once



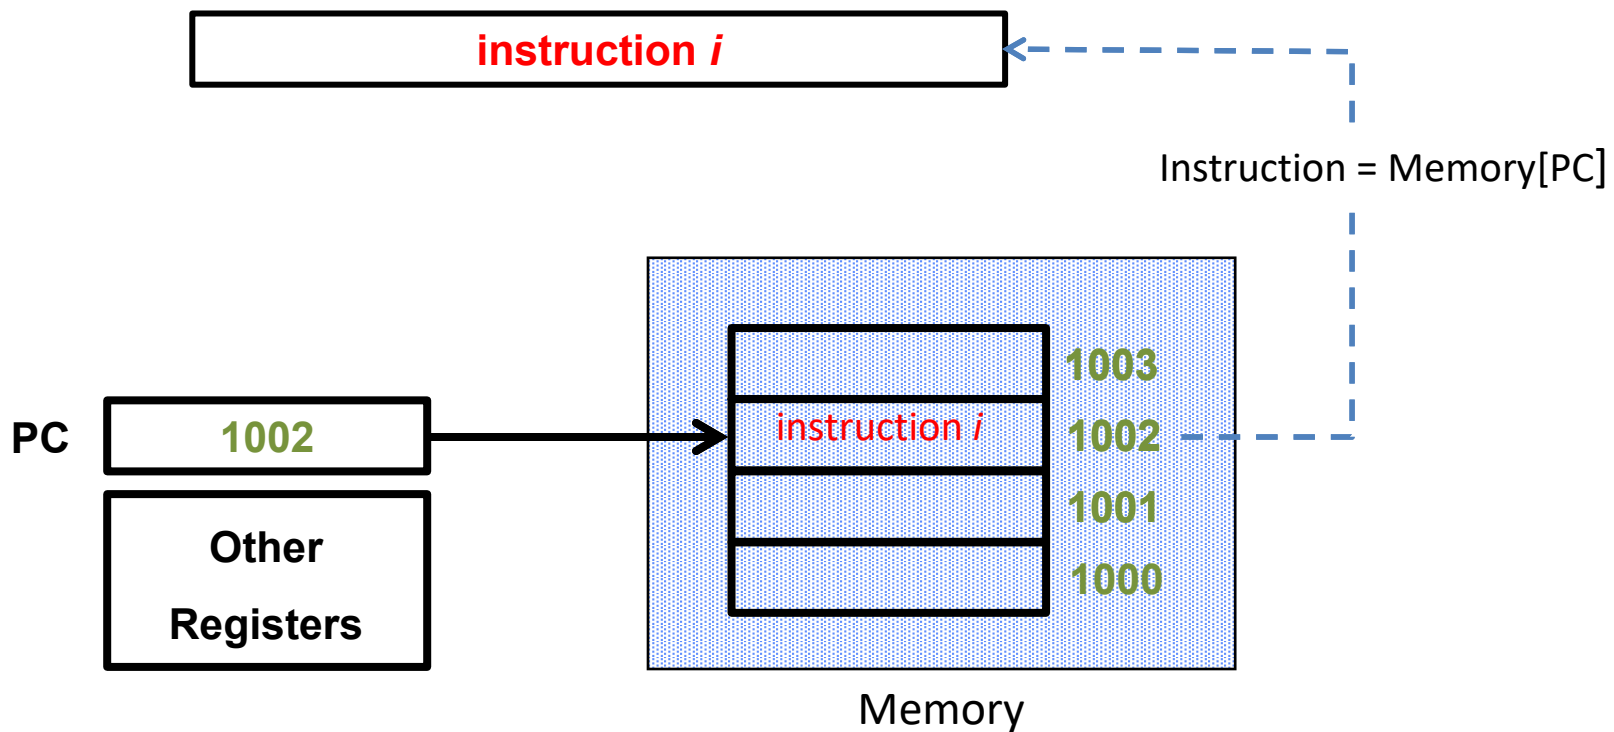
Instruction Granularity

- A statement such as $C = A + B$ is implemented as a sequence of primitive machine instructions that access memory at most once



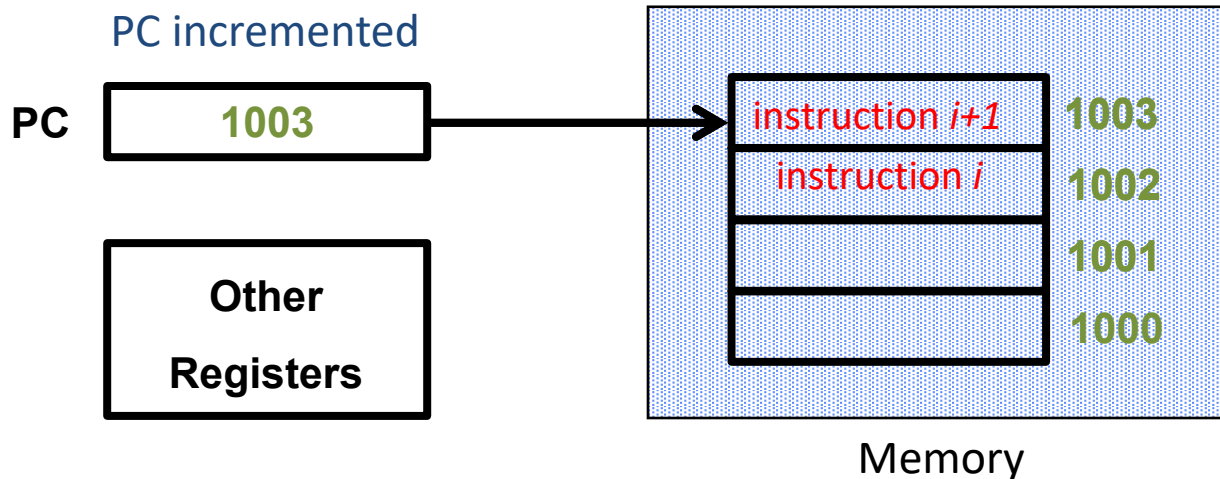
Program Counter

- PC is a register inside the processor
- PC holds the address of the next instruction to be fetched and executed



Program Counter

- PC is a register inside the processor
- PC holds the address of the next instruction to be fetched and executed
- After instruction is automatically read from memory, the PC is incremented to point to next instruction in memory



Instruction Cycle

- Set PC = address of first instruction in program
- Repeat
 - Fetch the instruction
 - Get an instruction from memory location contained in PC
 - Increment PC by size of instruction
 - Decode the instruction
 - Identify operation defined by instruction
 - If instruction requires data from memory, fetch data from memory
 - Execute the instruction
 - Perform operation defined by instruction
 - If instruction requires data to be stored, store data in memory
- Until (last instruction encountered)

Accessing the Program Counter

- The contents of the PC cannot be directly accessed by programs
- The content is manipulated by
 - branch, jump, and jump-to-subroutine instructions
 - exceptions (traps, interrupts)

Summary

- Instruction-Set Architecture (ISA)
 - Long-lived, stable interface between HW and SW
- Memory
 - Holds program instructions and data
 - Slow, large capacity, one-dimensional array
- Registers
 - Store values independent of memory
 - Small, but fast
- Program Counter
 - Holds address of next instruction to be read
 - Updated after instruction is read to point to next instruction
- Instruction Cycle
 - Fetch, decode, execute phases